

---

---

# Quantitative Analysis of KVM, Container, and Unikernel Environment Based on TCP/UDP Network Performance

**Bronco Loves Cloud**

Yinghe Chen Xi Wang Xiaoying Yang  
Miaoyan Zhang Mingjie Zhang

---

---

# Outline

- ❖ Introduction
- ❖ Environment Setup
- ❖ Experiment
- ❖ Result
- ❖ Conclusion

# Introduction

**KVM (Kernel-based Virtual Machine):** open source solution that converts Linux in a Type 1 hypervisor.

**Docker:** open source platform for the deployment and management of Linux Containers.

**Unikernel(OSv):** open source virtualization platform that run application directly on hardware / hypervisor.

# Introduction

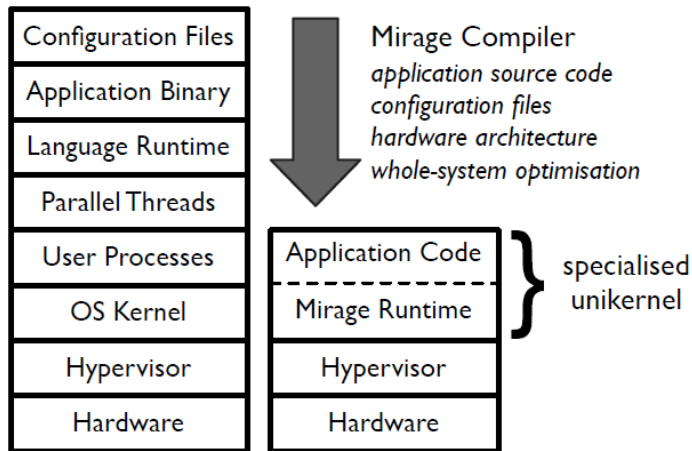


Figure 1: Contrasting software layers in existing VM appliances vs. unikernel's standalone kernel compilation approach.

## Unikernel vs KVM vs Docker

Less image size than docker and KVM

Less startup time / overhead than kvm

## Unikernel vs Linux-Network:

Throughput-higher

Transmit performance-lower


# Problem

**In this project, we quantitatively evaluate the network performance (TCP/UDP) for throughput using Netperf for different virtualization techniques: KVM, Container and Unikernel.**

# Environment Setup - Overall

Network benchmark: Netperf

One computer acts as the server, the other computer acts as the client, and the two computers are connected using ethernet directly.



Server Computer CPU	Intel(R) Pentium(R) CPU N3530 @ 2.16GHZ
Server Computer Memory	4G
Cable Speed	100Mbps

# Environment Setup - Host OS

Host OS: Ubuntu 14.04.5 LTS

To have a comparison, Netperf is installed on the host OS as well.

netperf-2.7.0, from <ftp://ftp.netperf.org/netperf/>

```
# Install netperf
wget ftp://ftp.netperf.org/netperf/netperf-2.7.0.tar.bz2
tar xf netperf-2.7.0.tar.bz2
cd netperf-2.7.0
./configure
make

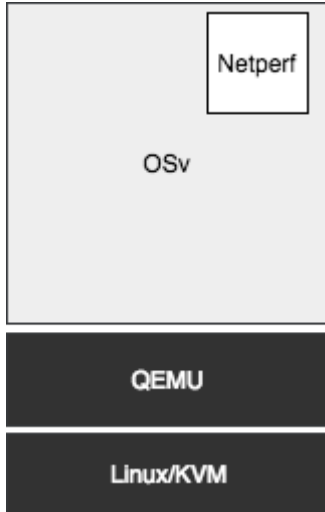
# Start netserver
cd netperf-2.7.0/src
./netserver -4
```

# Environment Setup - OSv Guest

The OSv installed on top of the hypervisor QEMU-KVM.

OSv image is built from source using source code from its official Git repository.

Netperf is cross-compiled on Linux host then incorporated into OSv image.



```
# Get OSv Source Code
git clone https://github.com/cloudius-systems/osv.git
```

```
# Install OSv Build Dependency
cd osv
python scripts/setup.py
git submodule update --init --recursive
```

```
# Build OSv Image
make
scripts/build
```

```
# Build netperf app
```

```
cd apps/netperf
make
```

```
# Include netserver in OSv image
cp apps/netperf/netserver.so tools/
```

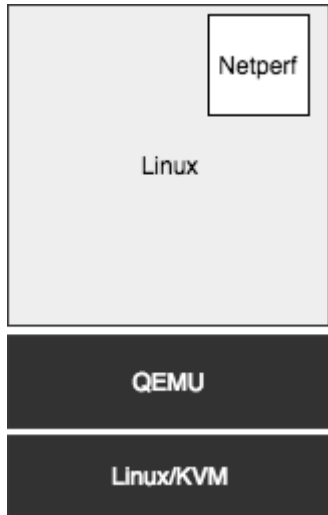
```
# Add `/tools/netserver.so: tools/mkfs/
netserver.so` to usr.manifest
scripts/build # Rebuild image
```

```
# Start netserver
sudo ./scripts/run.py -e "/tools/netserver.so -D
-4 -f" -c 4 --api
```



# Environment Setup - Ubuntu Guest

To test the network performance of KVM, we install a Linux guest on top of the hypervisor QEMU-KVM. Netperf is then installed on the Linux guest.



```
# Install QEMU KVM
sudo apt-get install qemu-kvm

# Create image
qemu-img create ubuntu.img 10G
```

```
# Download Ubuntu net install
http://archive.ubuntu.com/ubuntu/dists/trusty-updates/main/installer-amd64/current/images/netboot/mini.iso
```

```
# Start KVM guest with mini.iso as CD-ROM
qemu-system-x86_64 -hda ubuntu.img -cdrom
mini.iso -net nic -net user
```

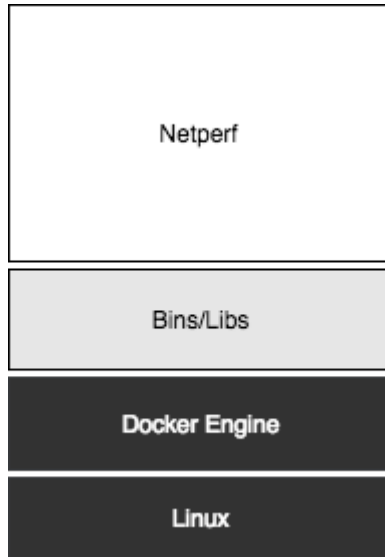
```
# Install Ubuntu as usual
# ...
```

```
# Start KVM guest again without CD-ROM
qemu-system-x86_64 -hda ubuntu.img -net
nic -net user
# Install netperf inside guest Ubuntu as usual
```

# Environment Setup - Docker

Docker is installed on the host Linux.

Netperf is installed in a docker container.



# Environment Setup - Network Configuration

Netperf server: one control port of TCP 12865 and one TCP or UDP data port. Ports of TCP 12865, TCP 12866 and UDP 12866 be mapped between host physical network and hypervisor virtual network.

```
# Start KVM guest with port mapping
# TCP 12865 is Netperf control channel
# TCP 12866 is used as Netperf TCP data channel, use -P 12866 to specify
# UDP 12866 is used as Netperf UDP data channel, use -P 12866 to specify
qemu-system-x86_64 -hda ubuntu.img -net nic -net user -redir tcp:12865::12865 -redir tcp:12866::12866 -redir udp:12866::12866

# For OSv, use the same -redir options by changing scripts/run.py
# Inside scripts/run.py change `args += ["-redir", "tcp:8000::8000"]`
# To `args += ["-redir", "tcp:8000::8000", "-redir", "tcp:12865::12865", "-redir", "tcp:12866::12866", "-redir", "udp:12866::12866"]`
sudo ./scripts/run.py -e "/tools/netserver.so -D -4 -f" -c 4 --api
```

# Environment Setup

All in all, we have the server with these software:

Host OS	Ubuntu 14.04.5 LTS
QEMU_KVM	Libvert 1.2.2 API: QEMU 1.2.2 Hypervisor: QEMU 2.0.0
OSv guest	v0.24
Ubuntu guest	Ubuntu 14.04.5 LTS
Docker	Docker 17.03.0-ce
Netperf	Netperf-2.7.0

# Netperf Experiment: TCP\_Stream Test

The TCP\_Stream test is to transfer some quantity of data from the system running netperf to the system running netserver.

```
yinghes-MacBook-Pro:src YC$ ./netserver
Starting netserver with host 'IN(6)ADDR_ANY' port '12865' and family AF_UNSPEC

yinghes-MacBook-Pro:src YC$ ./netperf
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 0 AF_INET to localhost ()
rt 0 AF_INET
Recv    Send      Send
Socket  Socket  Message  Elapsed
Size    Size    Size      Time    Throughput
bytes   bytes   bytes     secs.   10^6bits/sec

131072 131072 131072    10.00   34787.08
```

# Netperf Experiment: TCP\_Stream Test

```
$ /usr/etc/net_perf/netperf
TCP STREAM TEST ← the test type
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time  Throughput
bytes bytes bytes secs.  KBytes/sec

4096  4096  4096  10.00  2847.18 ← the performance
  ↑      ↑
  |      |
  |      | send size on the local system
  |      | receive size on the remote system

$ /usr/etc/net_perf/netperf -S 4096 -s 2048 -m 64 -H hwindfd9
TCP STREAM TEST
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time  Throughput
bytes bytes bytes secs.  KBytes/sec

4096  2048  64  10.00  18.05
  ↑      ↑      ↑
  |      |      |
  |      |      | changed by -s changed by -m
  |      |      |
  |      |      | changed by -S
  |      |      |
  |      |      | test with this system
```

# Netperf Experiment: TCP\_Stream Test

Control Variables:

1. IPV4 Vs. IPV6 (Internet Protocol Version) : IPV4
2. Port Selection : 12866
3. Message Size to transfer : 131072 bytes

Manipulated Variables:

TCP Socket Size (Start with 100 up to 20000 bytes)

# Netperf Experiment: TCP\_Stream Test

We chose python as our script language to run this experiment.

```
import subprocess
import time

f = open('TCP_Local.txt', 'a+')
for i in xrange(100, 131072, 1000):
    cmd = './netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s %s -m 131072' % i
    print(cmd)
    result = subprocess.check_output(cmd, shell=True)
    f.write(cmd + '\n')
    f.write(result)
    f.flush()
    time.sleep(60)
f.close()
```



# Netperf Experiment: TCP\_Stream Test

## OSV Test Result:

---

```
./netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s 100 -m 131072
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866
AF_INET
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time  Throughput
bytes bytes bytes secs.  10^6bits/sec

 65536   100 131072   10.00    4.85
./netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s 100 -m 131072
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866
AF_INET
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time  Throughput
bytes bytes bytes secs.  10^6bits/sec

 65536   100 131072   10.00    4.90
```

# Netperf Experiment: TCP\_Stream Test

## Docker Test Result

```
./netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s 100 -m 131072
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866
AF_INET
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time   Throughput
bytes bytes bytes secs.  10^6bits/sec

 87380   100 131072   10.00    3.85
./netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s 1100 -m 131072
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866
AF_INET
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time   Throughput
bytes bytes bytes secs.  10^6bits/sec

 87380   1100 131072   10.00    0.28
```

# Netperf Experiment: TCP\_Stream Test

## KVM Test Result

```
./netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s 100 -m 131072  
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866  
AF_INET
```

Recv Socket Size bytes	Send Socket Size bytes	Send Message Size bytes	Elapsed Time secs.	Throughput $10^6$ bits/sec
---------------------------------	---------------------------------	----------------------------------	--------------------------	-------------------------------

87380	100	131072	10.00	2.99
-------	-----	--------	-------	------

```
./netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s 1100 -m 131072  
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866  
AF_INET
```

Recv Socket Size bytes	Send Socket Size bytes	Send Message Size bytes	Elapsed Time secs.	Throughput $10^6$ bits/sec
---------------------------------	---------------------------------	----------------------------------	--------------------------	-------------------------------

87380	1100	131072	10.01	0.28
-------	------	--------	-------	------

# Netperf Experiment: TCP\_Stream Test

WE Find THE DIFFERENCE !!!

Recv Socket Size bytes	Send Socket Size bytes	Send Message Size bytes	Elapsed Time secs.	Throughput 10 <sup>6</sup> bits/sec
65536	100	131072	10.00	4.85
87380	100	131072	10.00	3.85
87380	100	131072	10.00	2.99

# Netperf Experiment: TCP\_Stream Test

What we really care is the throughput speed because it changes as the socket size changes

In order to prevent our experiment result is accidental, we runned second experiment with socket size starting from 300 to 20000 bytes with increase of 1000 each time.

AND WE FIND THE SAME THING !!!!

# But How Can We Compare the Performance ?

We use local OS as our benchmark and compare the ratio of specific environmental throughput to local OS throughput.

Local OS


```
./netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s 100 -m 131072
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866
AF_INET
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time  Throughput
bytes bytes bytes secs.  10^6bits/sec

87380 100 131072 10.00 4.33
./netperf -t TCP_STREAM -H 192.168.0.1 -- -P 12866 -s 1100 -m 131072
MIGRATED TCP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866
AF_INET
Recv  Send  Send
Socket Socket Message Elapsed
Size  Size  Size  Time  Throughput
bytes bytes bytes secs.  10^6bits/sec

87380 1100 131072 10.00 0.28
```

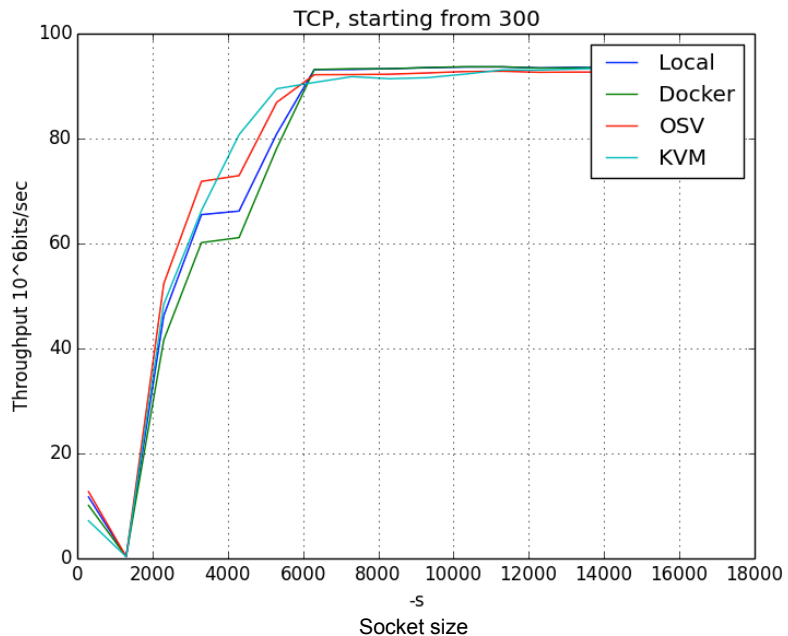
# Result

Netperf result: socket size starting from 100 bytes, with increment of 1000 bytes



# Result

Netperf result: socket size starting from 300 bytes, with increment of 1000bytes






# Result

Netperf result: combined result of TCP starting from 100 and 300 bytes

Normalized with respect to  
local running result




# Result

Netperf result: socket size starting from 300 bytes, with increment of 1000bytes


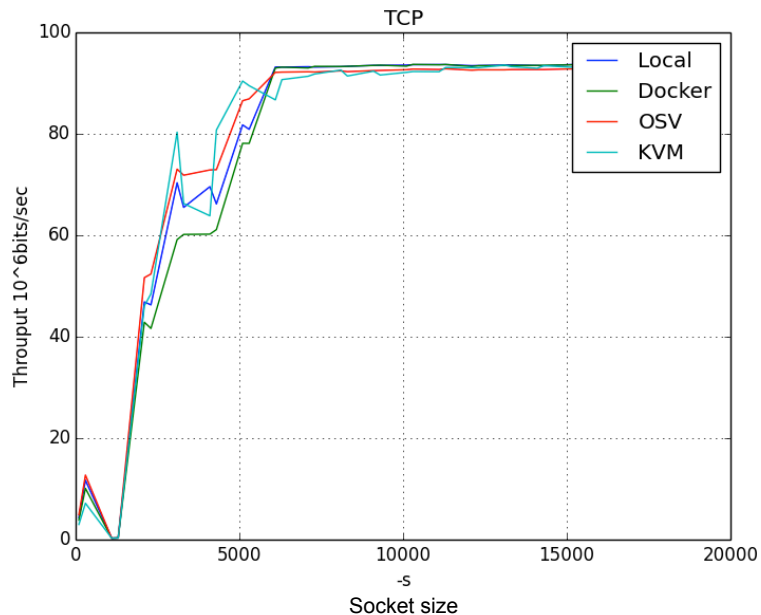
Normalized with respect to

local running result



# Result

Combination of result TCP running test



# Netperf Experiment: UDP\_Stream Test

Similar to the TCP\_Stream, we select -r requested size in our test as our manipulated variable, and python script to run this experiment.

```
import subprocess
import time

f = open('UDP_Local.txt', 'a+')
for i in xrange(1050, 3100, 100):
    cmd = './netperf -t UDP_STREAM -H 192.168.0.1 -- -P 12866 -r %s -m 3100' % i
    print(cmd)
    result = subprocess.check_output(cmd, shell=True)
    f.write(cmd + '\n')
    f.write(result)
    f.flush()
    time.sleep(60)
f.close()
```

# Netperf Experiment: UDP\_Stream Test

```
./netperf -t UDP_STREAM -H 192.168.0.1 -- -P 12866 -r 50 -m 3100  
MIGRATED UDP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866  
AF_INET
```

Socket Size bytes	Message Size bytes	Elapsed Time secs	Messages Okay Errors # #		Throughput 10^6bits/sec
9216	3100	10.00	1054894	89239	42.20
41600		10.00	321394		12.86

```
./netperf -t UDP_STREAM -H 192.168.0.1 -- -P 12866 -r 150 -m 3100  
MIGRATED UDP STREAM TEST from (null) (0.0.0.0) port 12866 AF_INET to (null) () port 12866  
AF_INET
```

Socket Size bytes	Message Size bytes	Elapsed Time secs	Messages Okay Errors # #		Throughput 10^6bits/sec
9216	3100	10.00	580762	651828	69.69
41600		10.00	316060		37.93

# Netperf Experiment: UDP\_Stream Test

```
$ /usr/etc/net_perf/netperf -t UDP_STREAM
UDP UNIDIRECTIONAL SEND TEST  ← the test type
Socket  Message  Elapsed  Messages
Size    Size    Time      Okay Errors  Throughput
bytes   bytes   secs      #         #         KBytes/sec

9216    9216    10.00     3413      0        3371.54 ← send performance
9360    9360    10.00     3334      0        3200.54

    ← send size on the local system    ← successful calls to send
```


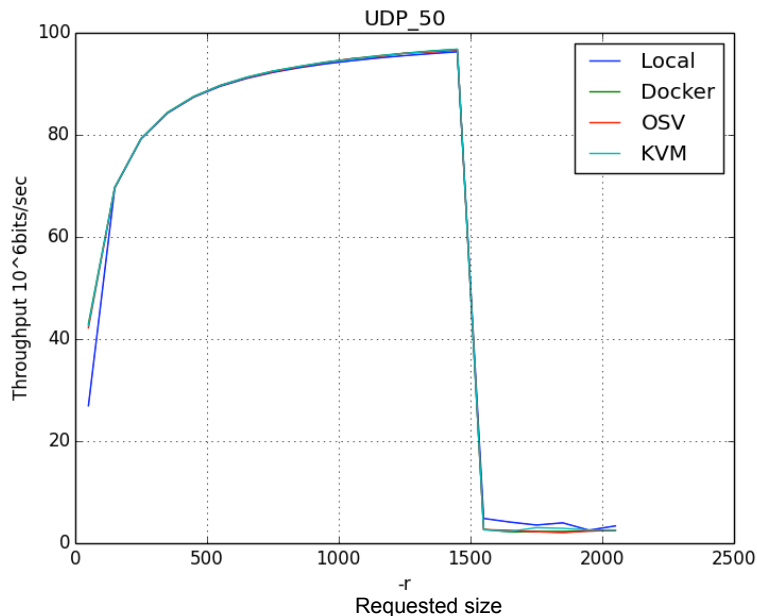
```
$ /usr/etc/net_perf/netperf -t UDP_STREAM -f m ← display megabits
UDP UNIDIRECTIONAL SEND TEST
Socket  Message  Elapsed  Messages
Size    Size    Time      Okay Errors  Throughput
bytes   bytes   secs      #         #         10^6bits/sec

9216    9216    10.00     3392      0        23.36
9360    9360    10.00     3374      0        23.73 ← receive perf

    ← receive size on the remote system    ← successful calls to recv
```

# Netperf Experiment: UDP\_Stream Test Result

Netperf result: requested size starting from 50 bytes, with increment of 100 bytes




# Netperf Experiment: UDP Error Comparison

It seems like the OSV has better

Okay rate in range of 250 to 900.

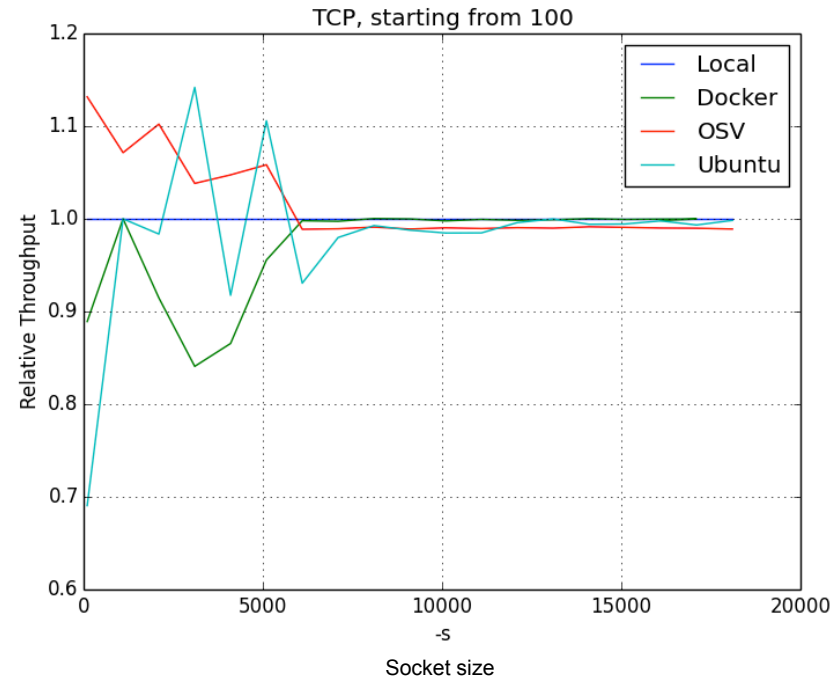
Okay rate =  $\# \text{ okay} / \#(\text{error} + \text{okay})$





# Conclusion

We calculate the performance based on  
The area that curve covered. The area  
Above Local OS is considered as positive, and  
Otherwise, negative. The performance is  
Determined by the total area.




# Conclusion

Thus under socket size restriction (0 - 5000), the TCP stream performance ranking is OSV, Ubuntu(KVM), and Docker

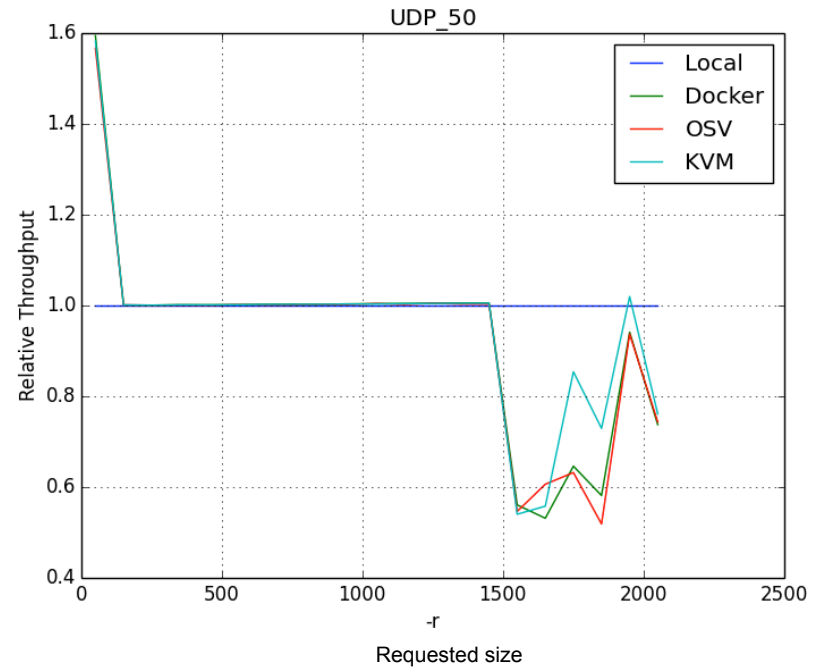
That unikernel outperformed others in

This case.



# Conclusion

Under request size restriction(1500 to 2000), KVM outperforms others in term of UDP stream performance.



# Contribution and Summary

We studied the three virtualization methods and did quantitatively analysis of the network performance under OSv, KVM, and Docker environment;

Quantitatively confirm the outperformance of OSv as claimed only qualitatively in papers.

# Future work

In the data collection process, we incremented socket size with increment of 1000 bytes, and the result of network performance in some environment experience bumping results.

This could be reduced by collecting data point with smaller increments, which will take longer time to finish. If more time allowed, the research could be rerun with more socket size point, thus will have more accurate results.

**Questions ?**